

OS Final Project Report

Abirbhav Dutta: ad5548

Yashika Dhawan: yd2281

Note: Throughout the report when we say 'with cache', we mean the cached scenario, and when we say without cache, we mean the non-cached scenario.

Part 1 : Basics

Program written: *run.c*.

Execution command: `./run <filename> [-r|-w] <block_size> <block_count>`

Output: *xor of all 4-byte integers for the block count and size specified*

Part 2 : Measurement

Program written: *run2.c*

Execution command: `./run2 <filename> <block_size>`

Output: *BlockCount*

Implementation details:

We ran the *run2.c* program for different block sizes starting with a small block size (1 Byte) and kept multiplying by 2. The *run2.c* program calls the code in *part1* internally. On finding a 'reasonable' time i.e. between 3 to 15 seconds, we returned the block count.

We also ran this program in a loop and kept doubling the block size to find the ideal block count for each block size. Result is given below.

| Block Size | Block Count | Time in seconds |
|------------|-------------|-----------------|
| 1 | 4194304 | 3.086139 |
| 2 | 4194304 | 3.005538 |
| 4 | 4194304 | 3.063287 |
| 8 | 4194304 | 3.083581 |
| 16 | 4194304 | 3.083881 |
| 32 | 4194304 | 3.162794 |
| 64 | 4194304 | 3.244777 |
| 128 | 4194304 | 3.444072 |

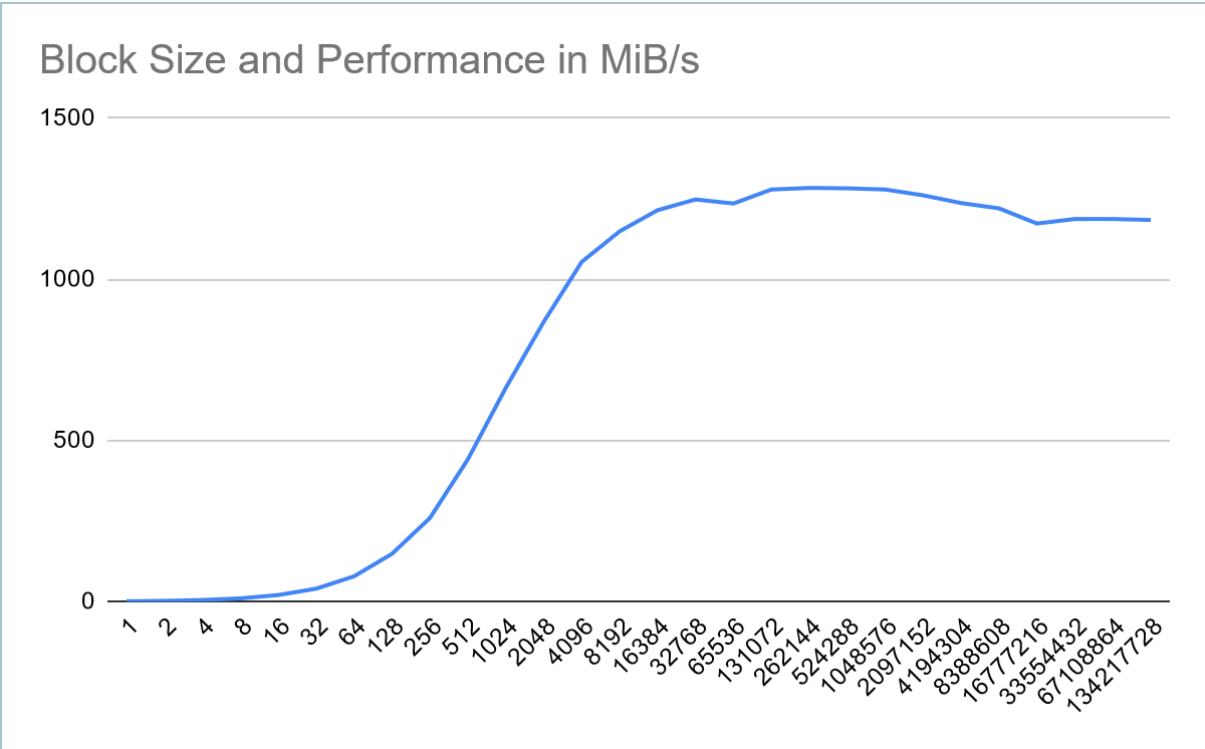
| | | |
|-----------|---------|----------|
| 256 | 4194304 | 3.936941 |
| 512 | 4194304 | 4.632057 |
| 1024 | 2097152 | 3.084435 |
| 2048 | 2097152 | 4.714712 |
| 4096 | 1048576 | 3.884537 |
| 8192 | 524288 | 3.563714 |
| 16384 | 262144 | 3.371497 |
| 32768 | 131072 | 3.281545 |
| 65536 | 65536 | 3.314986 |
| 131072 | 32768 | 3.202419 |
| 262144 | 16384 | 3.19015 |
| 524288 | 8192 | 3.193201 |
| 1048576 | 4096 | 3.202269 |
| 2097152 | 2048 | 3.248159 |
| 4194304 | 1024 | 3.311235 |
| 8388608 | 512 | 3.355782 |
| 16777216 | 256 | 3.490498 |
| 33554432 | 128 | 3.449231 |
| 67108864 | 64 | 3.449196 |
| 134217728 | 32 | 3.457746 |

Part 3 : Raw Performance

Note: For parts 3 to 5 we wrote the program *complete.c*.

Execution command: `./complete <filename>`

The graph below plots performance in MiB/s v/s Block Size. We can observe from the graph that the performance starts plateauing at around a block size of 8192, and the maximum performance is achieved for a block size of 262144 which is 256 KiB.



The same data in the form of a table:

Performance Table :

| Block Size | Performance in MiB/s |
|------------|----------------------|
| 1 | 1.296118 |
| 2 | 2.661753 |
| 4 | 5.223148 |
| 8 | 10.377544 |
| 16 | 20.75307 |
| 32 | 40.47055 |
| 64 | 78.896031 |
| 128 | 148.661225 |
| 256 | 260.100403 |
| 512 | 442.136149 |
| 1024 | 663.97905 |
| 2048 | 868.76993 |
| 4096 | 1054.43721 |
| 8192 | 1149.362604 |

| | |
|-----------|-------------|
| 16384 | 1214.890712 |
| 32768 | 1248.192421 |
| 65536 | 1235.60088 |
| 131072 | 1279.033273 |
| 262144 | 1283.952284 |
| 524288 | 1282.725269 |
| 1048576 | 1279.093044 |
| 2097152 | 1261.022008 |
| 4194304 | 1237.000596 |
| 8388608 | 1220.579755 |
| 16777216 | 1173.471526 |
| 33554432 | 1187.511185 |
| 67108864 | 1187.523005 |
| 134217728 | 1184.586712 |

Part 4 : Caching

Note: For parts 3 to 5 we wrote the program *complete.c*.

Execution command: `./complete <filename>`


Implementation details:

We ran the same tests above for each block size after clearing the cache.

To clear the cache, we used the command:

```
sudo sh -c "/usr/bin/echo 3 > /proc/sys/vm/drop_caches"
```

The result we got is shown below. The graph plots performance in MiB/s v/s Block Size:
Each test is run 5 times.




The same data in the form of a table:

| Block Size | Speed in MiB/s (Non-Cached) |
|------------|-----------------------------|
| 1 | 1.248193 |
| 2 | 2.504868 |
| 4 | 4.771619 |
| 8 | 9.95553 |
| 16 | 19.108872 |
| 32 | 36.058154 |
| 64 | 63.960244 |
| 128 | 98.434695 |
| 256 | 102.202593 |
| 512 | 111.358206 |
| 1024 | 113.583745 |
| 2048 | 116.353843 |
| 4096 | 119.854252 |
| 8192 | 120.407951 |
| 16384 | 115.946804 |
| 32768 | 119.509101 |

| | |
|-----------|------------|
| 65536 | 117.734969 |
| 131072 | 117.544451 |
| 262144 | 120.379179 |
| 524288 | 120.677978 |
| 1048576 | 117.754905 |
| 2097152 | 120.075822 |
| 4194304 | 99.23488 |
| 8388608 | 87.460922 |
| 16777216 | 91.233011 |
| 33554432 | 101.063456 |
| 67108864 | 105.893167 |
| 134217728 | 108.400913 |

What would be more interesting is if we plot both the cached and non-cached performances in the same graph. The following graph contains just that. This will help in comparing the performance between cached and non-cached reads. We can clearly see that cached reads are much faster.




Extra Credit: Comparing with dd


Implementation details: The code to run this experiment is also present in complete.c.

We ran the dd command for the same block sizes and block counts that we used in Part 3 and 4, and measured both the cached and non-cached performance. Comparison with our program is shown below. (NOTE that this is the comparison with our run.c file where dd outperforms run.c. However, as mentioned in Part 6, **our fast.c program outperforms dd**)

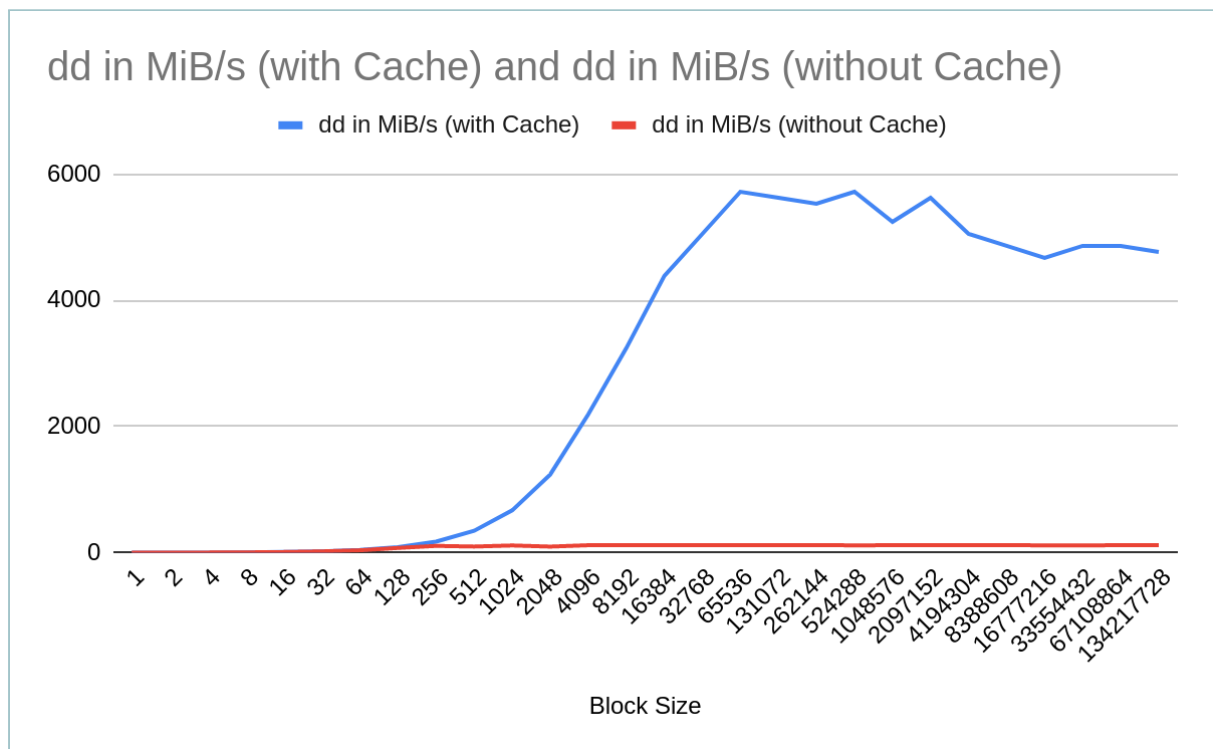
Below graph shows the comparison between the speed (performance) of our code with cache v/s speed of dd with cache.



Below graph shows the comparison between speed of our code without cache vs speed of dd without cache. (ie the non-cached performance).



Below graph shows the comparison of speed of dd with cache and without cache.



The same data in the form of a table:

| Block Size | Run.c speed in MiB/s (With Cache) | Run.c speed in MiB/s (Without Cache) | dd in MiB/s (with Cache) | dd in MiB/s (without Cache) |
|------------|-----------------------------------|--------------------------------------|--------------------------|-----------------------------|
| 1 | 1.296118 | 1.248193 | 1.236074038 | 1.190369211 |
| 2 | 2.661753 | 2.504868 | 2.538444631 | 2.388827485 |
| 4 | 5.223148 | 4.771619 | 4.981180446 | 4.550568978 |
| 8 | 10.377544 | 9.95553 | 9.896793897 | 9.494330117 |
| 16 | 20.75307 | 19.108872 | 19.79166328 | 18.2236344 |
| 32 | 40.47055 | 36.058154 | 38.5957113 | 34.38772396 |
| 64 | 78.896031 | 63.960244 | 75.24109347 | 60.99722174 |
| 128 | 148.661225 | 98.434695 | 141.7743451 | 93.87460932 |
| 256 | 260.100403 | 102.202593 | 248.0509917 | 97.46795568 |
| 512 | 442.136149 | 111.358206 | 421.6537498 | 106.1994257 |
| 1024 | 663.97905 | 113.583745 | 633.2195565 | 108.3218644 |
| 2048 | 868.76993 | 116.353843 | 828.5232942 | 110.9636349 |
| 4096 | 1054.43721 | 119.854252 | 1005.589352 | 114.3018839 |
| 8192 | 1149.362604 | 120.407951 | 1096.117232 | 114.8299323 |
| 16384 | 1214.890712 | 115.946804 | 1158.609685 | 110.5754524 |
| 32768 | 1248.192421 | 119.509101 | 1190.368659 | 113.9727224 |
| 65536 | 1235.60088 | 117.734969 | 1178.360434 | 112.2807788 |
| 131072 | 1279.033273 | 117.544451 | 1219.780778 | 112.0990868 |
| 262144 | 1283.952284 | 120.379179 | 1224.47191 | 114.8024932 |
| 524288 | 1282.725269 | 120.677978 | 1223.301738 | 115.08745 |
| 1048576 | 1279.093044 | 117.754905 | 1219.83778 | 112.2997913 |
| 2097152 | 1261.022008 | 120.075822 | 1202.603902 | 114.5131895 |
| 4194304 | 1237.000596 | 99.23488 | 1179.695306 | 94.63772495 |
| 8388608 | 1220.579755 | 87.460922 | 1164.035177 | 83.40920733 |
| 16777216 | 1173.471526 | 91.233011 | 1119.109284 | 87.00655053 |
| 33554432 | 1187.511185 | 101.063456 | 1132.498542 | 96.38159034 |
| 67108864 | 1187.523005 | 105.893167 | 1132.509814 | 100.9875601 |
| 134217728 | 1184.586712 | 108.400913 | 1129.709548 | 103.3791323 |

Extra Credit: Why '3'? In the command to clear cache?

Writing to *drop_cache* cleans cache without killing any application/service. Command *echo* is doing the job of writing to file.

echo 1 clears pagecache only. Pagecache is cached files; recently accessed files are stored here.

echo 2 clears dentries and inodes only. Dentries and inode cache are directory and file attributes.

echo 3 clears pagecache, inodes and dentries, i.e. all three.

Part 5 : System Calls

Note: For parts 3 to 5 we wrote the program *complete.c*.

Execution command: *./complete <filename>*

We got the following results for *read()* by changing the block size to 1 Byte:

Performance in MiB/s with Block Size of 1 byte : 1.296118 MiB/s

Performance in B/s with Block Size of 1 byte : 1359078.279162 B/s

We further tried with three system calls: *lseek*, *getpid* and *stat*. Each was run in a loop for a time greater than 5 seconds. The table below shows how *read* compares to *lseek*, *getpid* and *stat*.

| System Call | Speed (Sys Calls/sec) |
|---------------|-----------------------|
| <i>read</i> | 1359078.28 |
| <i>lseek</i> | 1870927.94 |
| <i>getpid</i> | 1976489.55 |
| <i>stat</i> | 971242.54 |

Part 6 : Raw Performance

Program written: *fast.c*.

Execution command: *./fast <filename>*

Output: xor of all 4-byte integers in the file.

We wrote two more programs here: *fasttest.c* and *fast2.c*. Details of the programs are mentioned below and the way to run is included in README.

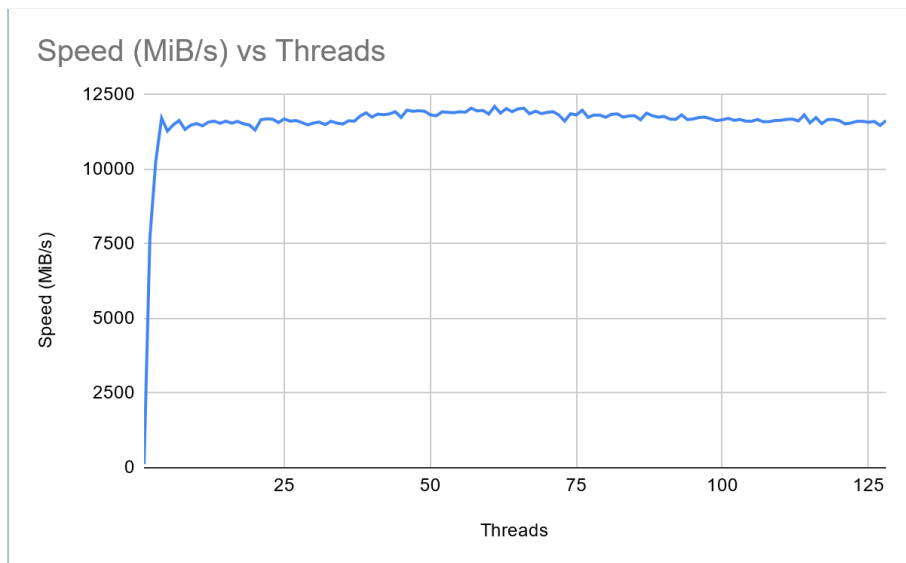
NOTE: Our program outperforms dd. On average, the speed for dd was 8.66 GiB/s, while the speed for our fast program is 11.18 GiB/s.

If you don't add block size (bs) to the dd command, the speed for dd is on average about 120.16 MB/s, which is even less.

| Fast.c | dd | dd (without bs) |
|-------------|------------|-----------------|
| 11.18 GiB/s | 8.66 GiB/s | 126 MB/s |

Details about how we implemented fast.c:

- We implemented multithreading, which significantly improved the performance.
- The block size that we chose was the block size for which we were getting the maximum value in Part 3, which turned out to be **256 KiB**.
- To select the ideal number of threads, we ran our program multiple times for different thread counts. The code for this is present in *fasttest.c*. You can find the graph below:



If we zoom in on the earlier threads, we get the following graph:




Table (Data for this graph): See Appendix (at the end)

The system we tested on had 4 logical processors, and we can clearly see from the graph that performance increases sharply till 4 threads. After that it mostly remains constant for larger threads. If we look at the data closely, we can see that performance keeps increasing slightly till about ~64 threads and then keeps decreasing slightly, but this change is a negligible amount, so not much can be concluded.

Cached v/s Non-Cached Performance (Fast.c)

To test our program's cached and non-cached performance, we fixed the threads to 64, and on an average got the following result:

| Cached (Fast.c) | Non-cached (Fast.c) |
|-----------------|---------------------|
| 11455.35 MiB/s | 91.51 MiB/s |

The fast.c program significantly outperforms run.c in the cached case. (Our performance for a block size of 256 KiB for run.c was around 1284 MiB/s). However, the non-cached result of fast.c is lesser than the result we got in Part 3 for our run.c program, where the speed was 120.379179 MiB/s for the block size of 256 KiB. To make sure we got the right result, we ran both run.c and fast.c in a **different system** after clearing the cache. For the block size of 256 KiB, we got the following result:

| Non-Cached (Fast.c) | Non-Cached (Run.c) |
|---------------------|--------------------|
| 2206.66 MiB/s | 383.19 MiB/s |

This result shows that our fast.c outperforms our run.c program in non-cached scenarios as well. Note that this result is contradictory to the performance in our own system.

Below is our thought process for choosing the number of threads in our final fast.c program:

Based on our observations we believe that taking a number of threads that is slightly more than the number of logical processors in a system seems a good and safe choice. We are doing the same in our file: *fast2.c* where we are taking the number of threads to be equal to 1.5 times the number of logical processors. However, there is no truly portable code to get the logical processors of a system. The solution we used: `int num_logical = sysconf(_SC_NPROCESSORS_CONF);` might not work on every system. Therefore we didn't do this in our fast.c program. (This code is present in fast2.c that can be checked).

The number of threads that we chose finally was **128**. The m5d.metal instance has 96 logical processors, and therefore we chose our number of threads to be greater than that number, assuming that the system in which our code will be tested will be of a similar or lesser specification to m5d.metal.

Appendix

Speed (MiB/s) v/s Threads

| Threads | Speed (MiB/s) |
|---------|---------------|
| 1 | 120.209326 |
| 2 | 7712.755072 |
| 3 | 10267.98198 |
| 4 | 11717.65752 |
| 5 | 11282.45554 |
| 6 | 11492.7167 |
| 7 | 11637.24336 |
| 8 | 11339.41729 |
| 9 | 11481.52836 |

| | |
|----|-------------|
| 10 | 11532.30768 |
| 11 | 11457.88937 |
| 12 | 11581.30129 |
| 13 | 11616.28564 |
| 14 | 11543.05162 |
| 15 | 11615.17273 |
| 16 | 11547.72143 |
| 17 | 11608.29655 |
| 18 | 11525.54102 |
| 19 | 11490.01124 |
| 20 | 11319.13127 |
| 21 | 11663.25754 |
| 22 | 11693.18277 |
| 23 | 11681.96401 |
| 24 | 11568.58657 |
| 25 | 11691.00033 |
| 26 | 11614.01218 |
| 27 | 11634.67354 |
| 28 | 11571.26978 |
| 29 | 11490.65529 |
| 30 | 11554.31288 |
| 31 | 11583.39536 |
| 32 | 11498.11971 |
| 33 | 11611.27323 |
| 34 | 11550.53697 |
| 35 | 11517.7821 |
| 36 | 11628.28992 |
| 37 | 11613.07912 |
| 38 | 11788.31003 |
| 39 | 11894.13773 |
| 40 | 11753.10424 |
| 41 | 11852.33942 |
| 42 | 11832.10276 |
| 43 | 11857.41252 |
| 44 | 11929.66463 |

| | |
|----|-------------|
| 45 | 11740.42535 |
| 46 | 11984.01066 |
| 47 | 11946.62796 |
| 48 | 11965.23933 |
| 49 | 11948.69163 |
| 50 | 11828.30444 |
| 51 | 11792.81039 |
| 52 | 11925.27365 |
| 53 | 11911.31469 |
| 54 | 11898.93274 |
| 55 | 11928.86947 |
| 56 | 11914.80137 |
| 57 | 12050.89797 |
| 58 | 11958.05669 |
| 59 | 11975.36803 |
| 60 | 11853.69764 |
| 61 | 12108.48898 |
| 62 | 11883.25661 |
| 63 | 12036.862 |
| 64 | 11933.52833 |
| 65 | 12026.85996 |
| 66 | 12049.17227 |
| 67 | 11861.96581 |
| 68 | 11951.13603 |
| 69 | 11868.24634 |
| 70 | 11907.67914 |
| 71 | 11932.35381 |
| 72 | 11822.94626 |
| 73 | 11618.14097 |
| 74 | 11865.57349 |
| 75 | 11823.09504 |
| 76 | 11981.87112 |
| 77 | 11744.48549 |
| 78 | 11818.8317 |
| 79 | 11816.96127 |

| | |
|-----|-------------|
| 80 | 11746.12502 |
| 81 | 11843.5248 |
| 82 | 11861.17965 |
| 83 | 11750.6422 |
| 84 | 11790.78788 |
| 85 | 11796.26505 |
| 86 | 11661.37571 |
| 87 | 11883.30671 |
| 88 | 11799.79582 |
| 89 | 11748.23016 |
| 90 | 11777.25252 |
| 91 | 11684.55477 |
| 92 | 11674.08984 |
| 93 | 11825.79839 |
| 94 | 11667.50594 |
| 95 | 11688.05533 |
| 96 | 11738.11526 |
| 97 | 11752.47945 |
| 98 | 11694.51689 |
| 99 | 11633.83319 |
| 100 | 11663.60743 |
| 101 | 11707.85048 |
| 102 | 11643.25205 |
| 103 | 11672.2166 |
| 104 | 11614.11985 |
| 105 | 11611.21345 |
| 106 | 11671.70911 |
| 107 | 11594.28243 |
| 108 | 11599.24438 |
| 109 | 11636.5948 |
| 110 | 11643.04766 |
| 111 | 11676.2176 |
| 112 | 11684.5911 |
| 113 | 11621.19458 |
| 114 | 11823.51658 |

| | |
|-----|-------------|
| 115 | 11559.03921 |
| 116 | 11734.07173 |
| 117 | 11534.19528 |
| 118 | 11667.80779 |
| 119 | 11673.4855 |
| 120 | 11635.17781 |
| 121 | 11522.12534 |
| 122 | 11553.47226 |
| 123 | 11608.46387 |
| 124 | 11610.71125 |
| 125 | 11578.25672 |
| 126 | 11604.29418 |
| 127 | 11471.87898 |
| 128 | 11630.40102 |